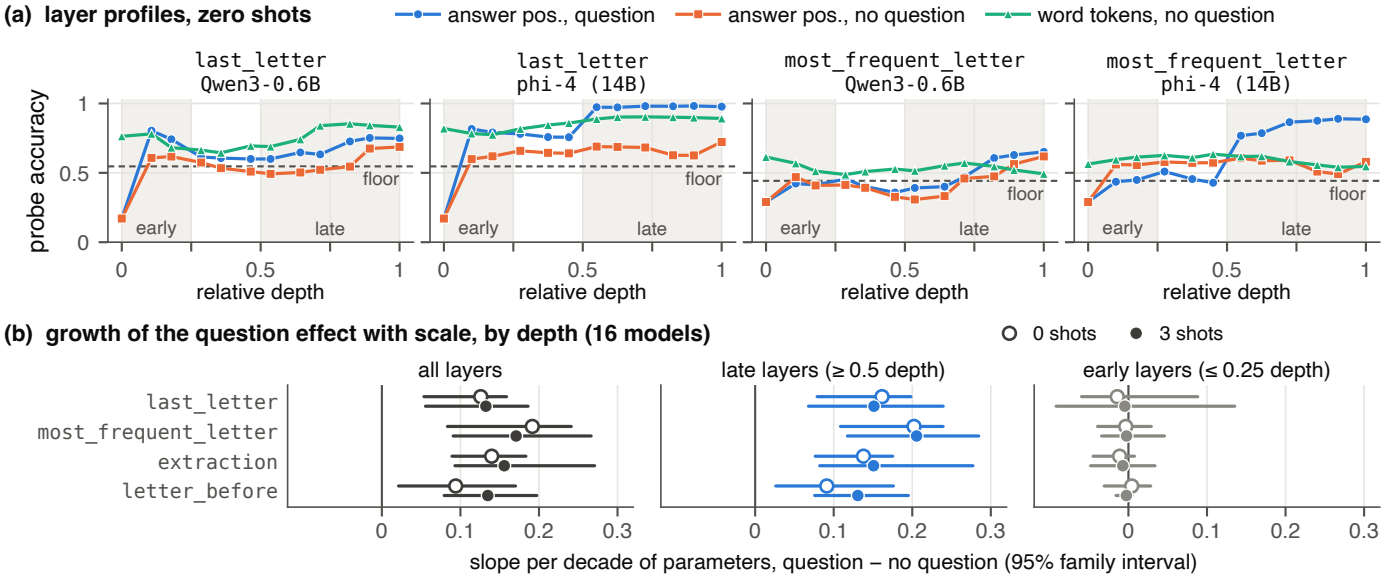
(a) layer profiles, zero shots
answer pos., question
answer pos., no question
word tokens, no question
last_letter
Qwen3-0.6B
last_letter
phi-4 (14B)
most_frequent_letter
Qwen3-0.6B
most_frequent_letter
phi-4 (14B)
probe accuracy
1
0.5
0
floor
early
late
0
0.5
1
relative depth
(b) growth of the question effect with scale, by depth (16 models)
0 shots
3 shots
all layers
late layers (≥ 0.5 depth)
early layers (≤ 0.25 depth)
last_letter
most_frequent_letter
extraction
letter_before
0
0.1
0.2
0.3
slope per decade of parameters, question – no question (95% family interval)